

# **Boosting Generalization in Diffusion-Based Neural Combinatorial Solver via Energy-guided Sampling**

---

윤성준

# Table of Contents

- I Introduction**
- II Related works**
- III Theoreitcal Analysis**
- IV Proposed Approach**
- V Numerical Results**

# How to solve Combinatorial Optimizaion Problems?


## Combinatorial Optimizaion Problems

### Exact Solver

- Optimal Solution
- Large computation Complexity

### Heuristic & Approximation

- Find Feasible Solution faster than exact solver
- Certain level of optimality




### Deep Learning



- Fast Compuation time
- Less Optimality gap

# I. Introduction

## 1. Cross-scale problem



## 2. Cross-problem transfer capability



# I. Introduction

## Adapter Foundation

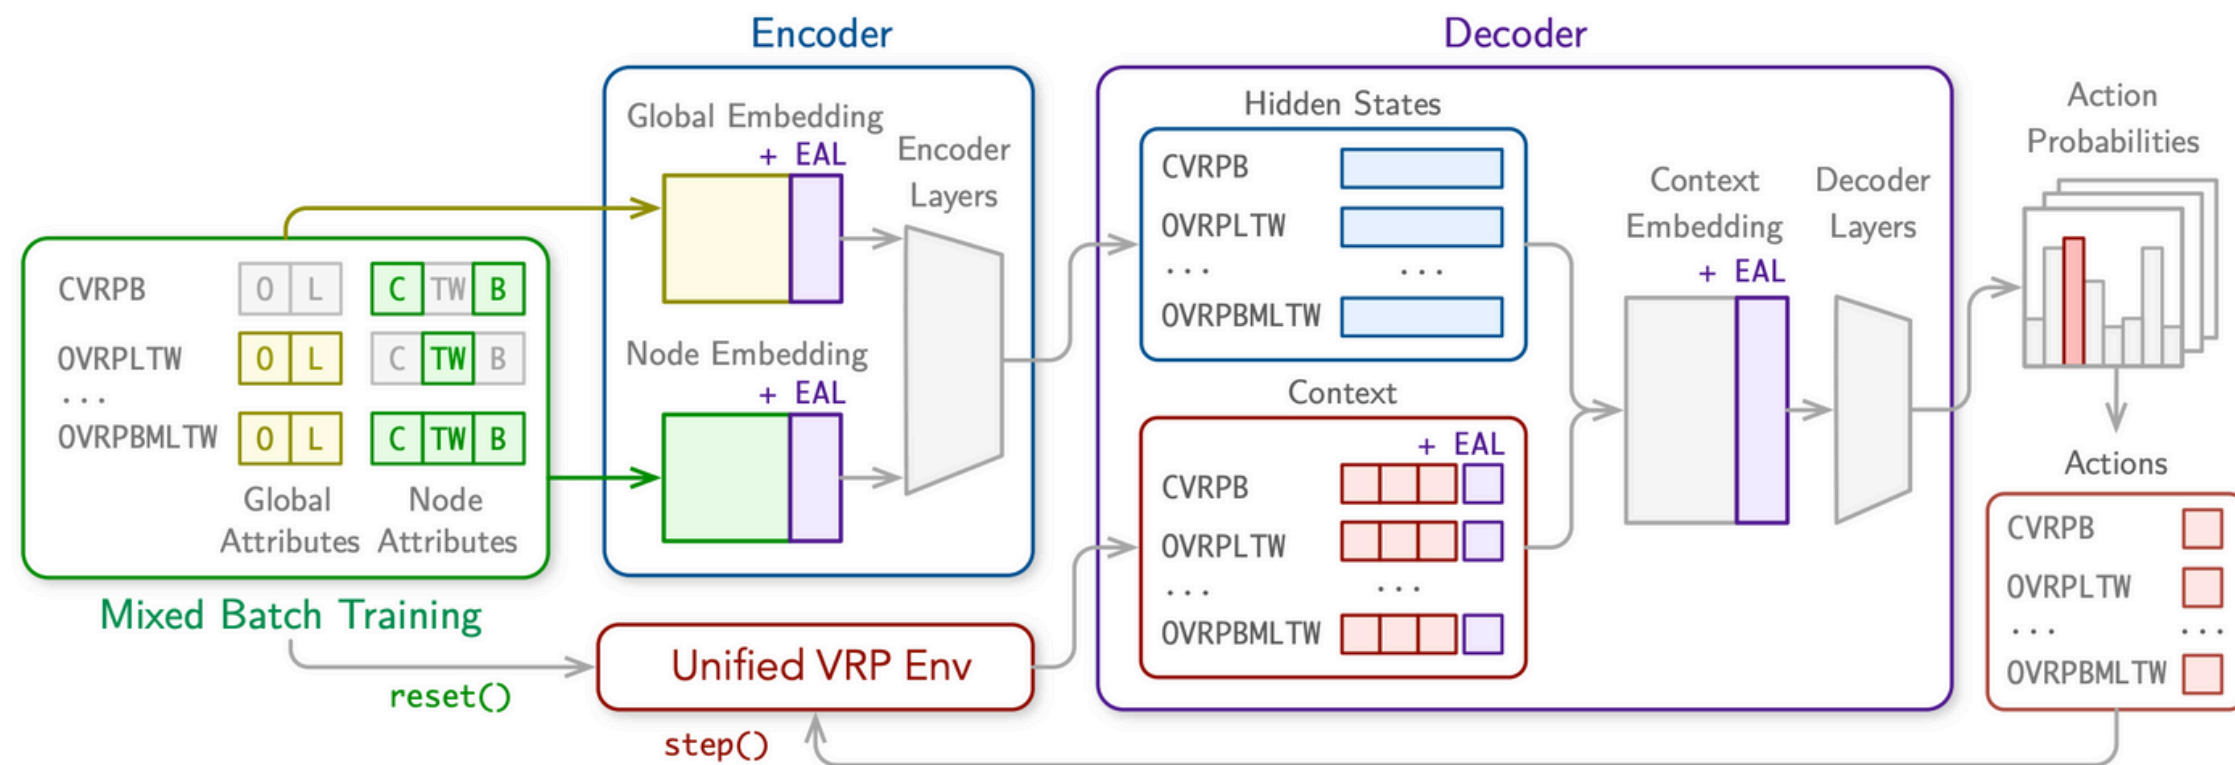


Figure 4.1: Overview of RouteFinder.

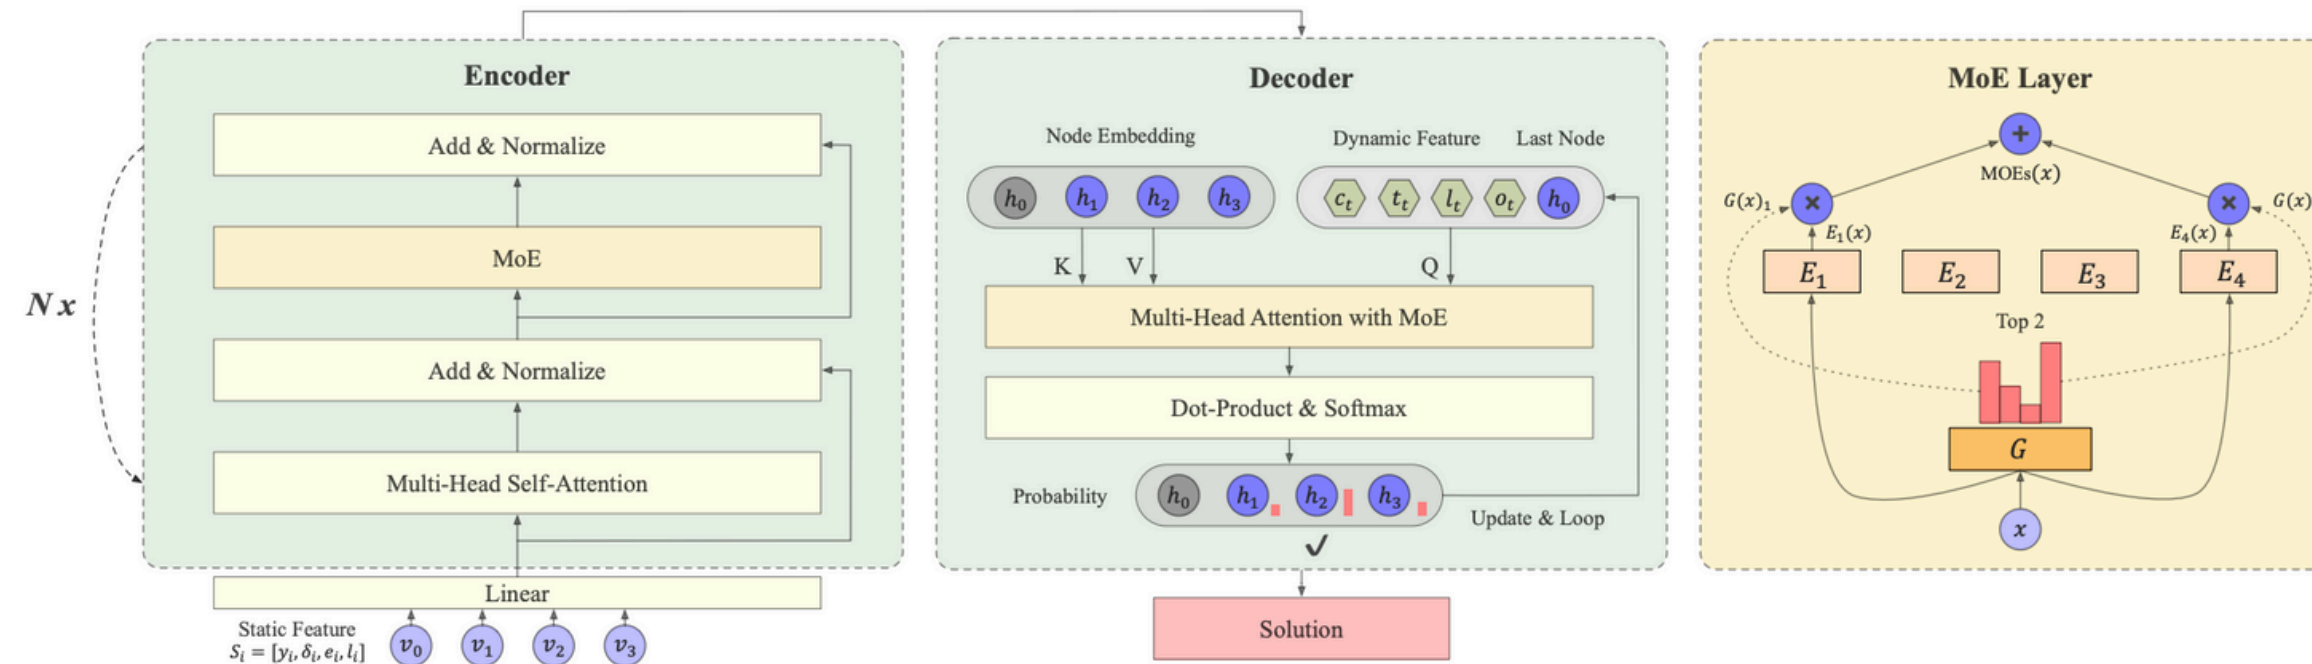
Reference:

RouteFinder Towards Foundation Models for Vehicle Routing Problems

<https://openreview.net/pdf?id=hCiaiZ6e4G>

# I. Introduction

## MOE Foundation



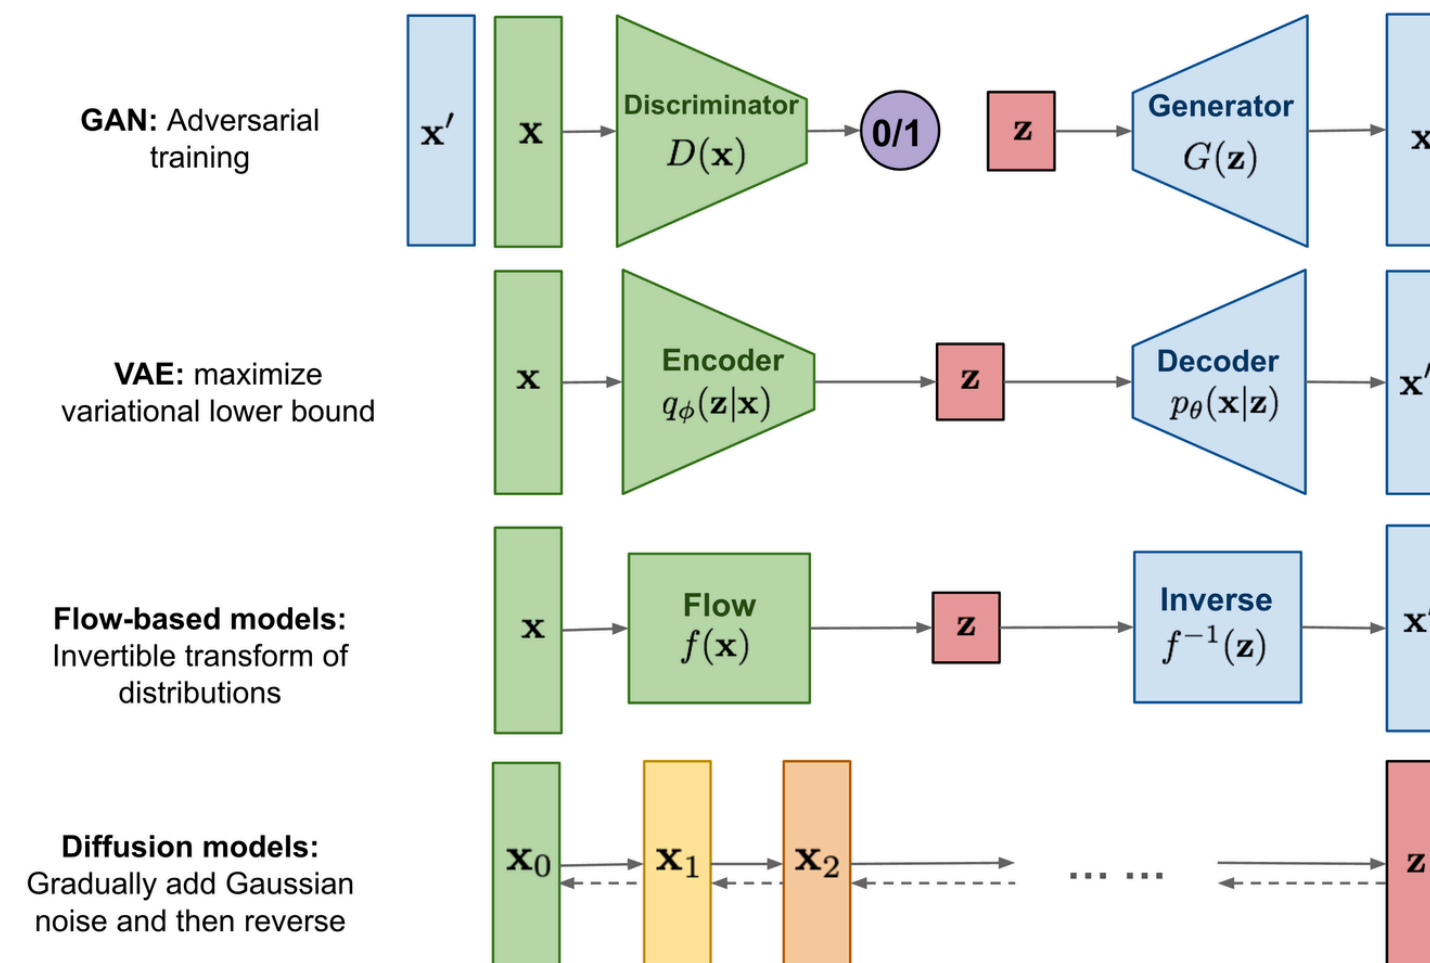
Reference:

MVMoE: Multi-Task Vehicle Routing Solver with Mixture-of-Experts

<https://arxiv.org/pdf/2405.01029>

## II. Related works

### Generative Model(Diffusion)



parameterized Markov chain trained using variational inference to produce samples matching the data after finite time

$$p_\theta(X_0) := \int p_\theta(X_{0:T}) dX_{1:T}$$

## II. Related works

### Guided Sampling

### Classifier guidance

$$p_{\theta, \phi}(x_t | x_{t+1}, y) = Z p_{\theta}(x_t | x_{t+1}) p_{\phi}(y | x_t)$$

### Classifier-Free guidance

$$\tilde{\epsilon}_{\theta}(\mathbf{z}_{\lambda}, \mathbf{c}) = (1 + w) \epsilon_{\theta}(\mathbf{z}_{\lambda}, \mathbf{c}) - w \epsilon_{\theta}(\mathbf{z}_{\lambda})$$

### Training-Free guidance




Figure 2: The directed graphical model considered in this work.

Reference:

Diffusion Models Beat GANs on Image

<https://arxiv.org/pdf/2105.05233#page=25&zoom=100,144,96>

CLASSIFIER-FREE DIFFUSION GUIDANCE

<https://arxiv.org/pdf/2207.12598>

Universal Guidance for Diffusion Models

<https://arxiv.org/pdf/2302.07121>



### III. Theoretical Analysis

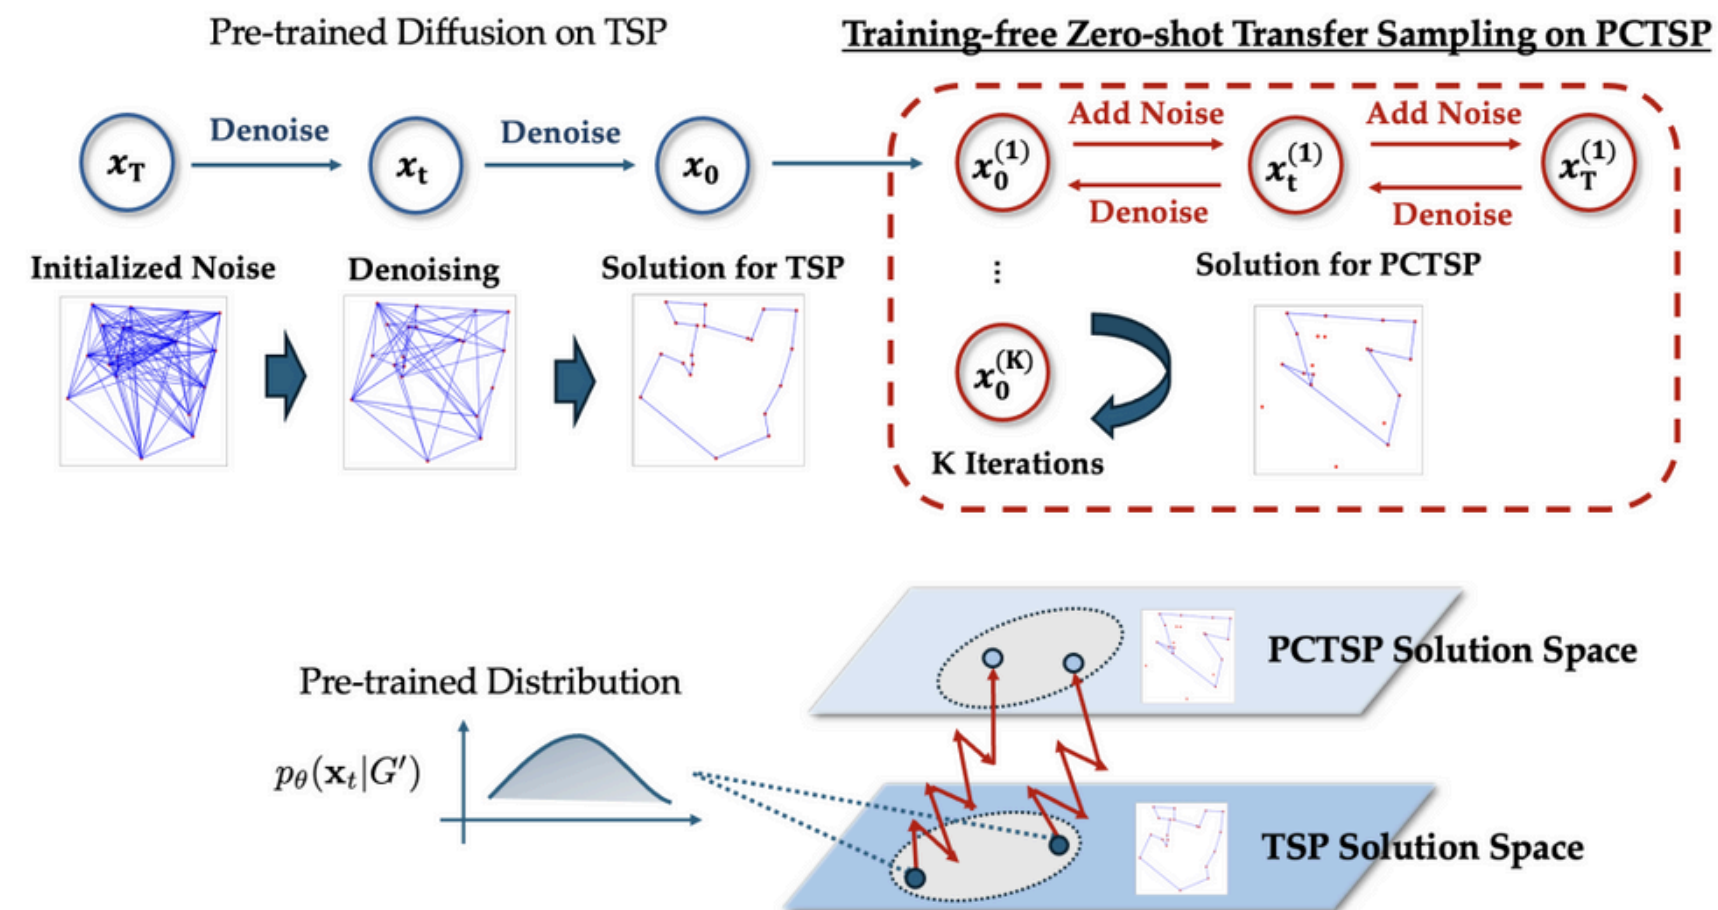




Figure 1: Overview of energy-guided sampling framework for achieving cross-problem generalization. Left: Pre-trained diffusion model performs denoising on original problem  $G$  (TSP). Right: Proposed energy-guided sampling on target problem  $G'$  (PCTSP) conducts  $K$  rewrite iterations of noise addition and guided denoising.

### III. Theoretical Analysis



Combinatorial Optimization problem



TSP




Prize Collecting TSP(PCTSP)



Orienteering Problem

### III. Theoretical Analysis

Combinatorial Optimization problem




Prize Collecting TSP(PCTSP)

The goal is to collect as much prize as possible while minimizing the total travel cost.  
The prize is only revealed when the node is visited.  
if not visit node, each node have penalty

### III. Theoretical Analysis

#### Combinatorial Optimization problem



Orienteering Problem

At each step, the agent chooses a location to visit in order to maximize the collected prize.  
The total length of the path must not exceed a given threshold.

### III. Theoretical Analysis

#### CO problem objective Formulation

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in \mathcal{X}_G} f(\mathbf{x}; G)$$

$$f(\mathbf{x}; G) = f_{\text{cost}}(\mathbf{x}; G) + \beta \cdot f_{\text{valid}}(\mathbf{x}; G)$$

$f_{\text{cost}}(\cdot; G)$  : measures the solution quality

$f_{\text{valid}}(\cdot; G)$  : problem- specific constraints through a penalty coefficient

$\mathbf{x} \in \{0, 1\}^N$  : decision variable

# III. Theoretical Analysis

## CO problem objective Formulation

**Proposition A.5** (Prize Collecting TSP (PCTSP)). *Given a complete graph  $G = (V, E)$  with edge weights  $w : E \rightarrow \mathbb{R}^+$ , vertex prizes  $r : V \rightarrow \mathbb{R}^+$ , penalties  $p : V \rightarrow \mathbb{R}^+$ , and prize threshold  $R$ , find  $\mathbf{x} \in \{0, 1\}^{|E|}$ ,  $\mathbf{y} \in \{0, 1\}^{|V|}$  that minimizes:*

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in \mathcal{X}_G} f(\mathbf{x}; G)$$

$$f(\mathbf{x}; G) = f_{\text{cost}}(\mathbf{x}; G) + \beta \cdot f_{\text{valid}}(\mathbf{x}; G)$$

$$f(\mathbf{x}, \mathbf{y}; G) = f_{\text{cost}}(\mathbf{x}, \mathbf{y}; G) + \beta \cdot f_{\text{valid}}(\mathbf{x}, \mathbf{y}; G)$$

$$\text{where } f_{\text{cost}}(\mathbf{x}, \mathbf{y}; G) = \sum_{e \in E} w_e x_e + \sum_{v \in V} p_v (1 - y_v)$$

$$f_{\text{valid}}(\mathbf{x}, \mathbf{y}; G) = \max(0, R - \sum_{v \in V} r_v y_v)$$

$f_{\text{cost}}(\cdot; G)$  : measures the solution quality

$f_{\text{valid}}(\cdot; G)$  : problem- specific constraints through a penalty coefficient

$\mathbf{x} \in \{0, 1\}^N$  : decision variable

# III. Theoretical Analysis

## CO problem objective Formulation

**Proposition A.6** (Orienteering Problem (OP)). *Given a complete graph  $G = (V, E)$  with edge weights  $w : E \rightarrow \mathbb{R}^+$ , vertex scores  $s : V \rightarrow \mathbb{R}^+$ , and budget  $B$ , find  $\mathbf{x} \in \{0, 1\}^{|E|}$ ,  $\mathbf{y} \in \{0, 1\}^{|V|}$  that minimizes:*

$$f(\mathbf{x}, \mathbf{y}; G) = f_{\text{cost}}(\mathbf{x}, \mathbf{y}; G) + \beta \cdot f_{\text{valid}}(\mathbf{x}, \mathbf{y}; G)$$

$$\text{where } f_{\text{cost}}(\mathbf{x}, \mathbf{y}; G) = - \sum_{v \in V} s_v y_v$$

$$f_{\text{valid}}(\mathbf{x}, \mathbf{y}; G) = \max(0, \sum_{e \in E} w_e x_e - B)$$

$f_{\text{cost}}(\cdot; G)$  : measures the solution quality

$f_{\text{valid}}(\cdot; G)$  : problem- specific constraints through a penalty coefficient

$\mathbf{x} \in \{0, 1\}^N$  : decision variable

### III. Theoretical Analysis

#### Analysis transfer of Combinatorial Optimization problem (PCTSP)

**Definition 4.1** (Marginal Decrease). For a non-empty subset of nodes  $S \subseteq V$ , let  $\text{TSP}(S)$  denote the cost of the optimal TSP tour visiting all nodes in  $S$ . The marginal decrease of a subset  $S$  is defined as

$$\Delta(S) = \text{TSP}(V) - \text{TSP}(V \setminus S).$$

The marginal decrease measures the cost reduction of not visiting a subset of nodes  $S$ , which helps quantify the difference between the optimal tours. Take PCTSP as an example. If for any non-empty subset of nodes  $S \subseteq V$ , the penalty of not visiting the nodes in  $S$  satisfies  $\sum_{i \in S} p_i \geq \Delta(S)$ , then PCTSP and TSP share the same optimal tours. Based on this notion, we formalize the structural similarities in the following theorem.

**TSP(V) : cost of the optimal TSP**

**TSP(V\S): cost of the optimal TSP except for subset S**

**\Delta(S): cost reduction TSP except for subset S**



### III. Theoretical Analysis

#### Analysis transfer of Combinatorial Optimization problem (PCTSP)

Under these assumptions, we compare the optimal solutions of PCTSP/OP with that of TSP as follows.

(i) For PCTSP, if for any non-empty subset of nodes  $S \subseteq V$ , the penalty of not visiting the nodes in  $S$  satisfies

$$\sum_{i \in S} p_i \geq \Delta(S),$$

then PCTSP and TSP share the same optimal solutions, since not visiting any subset of nodes would result in a higher total cost. We can thus formulate the optimal cost (tour length + penalty) of PCTSP as

$$\text{PCTSP}(V) = \text{TSP}(V) + \min_{S \subseteq V} \sum_{i \in S} p_i - \Delta(S). \quad (10)$$

And we have the optimal solutions of PCTSP being the tours in  $\text{TSP}(V \setminus S)$  with

$$S \in \arg \min_{S \subseteq V} \sum_{i \in S} p_i - \Delta(S),$$

completing the proof for PCTSP.

### III. Theoretical Analysis

#### Analysis transfer of Combinatorial Optimization problem




Figure 2: Illustration of policy optimization trajectories on VRP with varying difficulty levels - (a)(b)(d) easy and (c)(e) hard, and different constraint handling schemes - (a) feasibility masking, (b)(c) Lagrangian multiplier, and (d)(e) our PIP. The orange-filled circle denotes the feasible policy space  $\Pi_F$ , while the dotted frame represents the actual search space of the neural policies  $\pi_\theta$ .

Reference:

Learning to Handle Complex Constraints for Vehicle Routing Problems

<https://arxiv.org/pdf/2410.21066>

### III. Theoretical Analysis

Probabilistic Modeling for CO

$$p_{\theta, \phi}(x_t | x_{t+1}, y) = Z p_{\theta}(x_t | x_{t+1}) p_{\phi}(y | x_t)$$

$$\mathcal{E}(\cdot | \mathcal{G}) \doteq |y - f(\cdot; G)|$$

$$p(y | \mathbf{x}; G) = \frac{\exp\left(-\frac{1}{\tau} \mathcal{E}(y, \mathbf{x}; G)\right)}{\mathcal{Z}}, \text{ where } \mathcal{Z} = \sum_{\mathbf{x}} \exp\left(-\frac{1}{\tau} \mathcal{E}(y, \mathbf{x}; G)\right),$$

$\tau$ : controls the temperature of system

$Z$ : partition function for pmf condition

$\mathbf{x}$ : optimal solution

$$L(\theta) = \mathbb{E}_{G \sim \mathcal{G}}[-\log p_{\theta}(\mathbf{x} | y_G^*, G)].$$

# III. Theoretical Analysis

## Discrete Diffusion Generation Modeling

### Loss

$$\begin{aligned} L(\theta) &= \mathbb{E}_{G \sim \mathcal{G}}[-\log p_{\theta}(\mathbf{x}_0|G)] \\ &\leq \sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} \left[ D_{KL}[q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) || p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t, G)] - \log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1, G) \right] + C. \end{aligned}$$

### III. Theoretical Analysis

#### Discrete Diffusion Generation Modeling(DIFUSCO)

$$\begin{aligned} q(x_{1:T}|x_0) &= \prod_{t=1}^T q(x_t|x_{t-1}) & q(\mathbf{x}_t|\mathbf{x}_{t-1}) &= \text{Cat}(\mathbf{x}_t; \mathbf{p} = \tilde{\mathbf{x}}_{t-1} \mathbf{Q}_t), & \text{forward process} \\ p_\theta(x_{0:T}|G) &= p(x_T) \prod_{t=1}^T p_\theta(x_{t-1}|x_t, G) & p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, G) &= \sum_{\tilde{\mathbf{x}}_0} q(\mathbf{x}_{t-1}|\mathbf{x}_t, \tilde{\mathbf{x}}_0) p_\theta(\tilde{\mathbf{x}}_0|\mathbf{x}_t, G). & \text{backward process} \end{aligned}$$

$$\begin{aligned} q(\mathbf{x}_t|\mathbf{x}_0) &= \text{Cat}(\mathbf{x}_t; \mathbf{p} = \tilde{\mathbf{x}}_0 \overline{\mathbf{Q}}_t), \\ q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) &= \text{Cat} \left( \mathbf{x}_{t-1}; \mathbf{p} = \frac{\tilde{\mathbf{x}}_t \mathbf{Q}_t^\top \odot \tilde{\mathbf{x}}_0 \overline{\mathbf{Q}}_{t-1}}{\tilde{\mathbf{x}}_0 \overline{\mathbf{Q}}_t \tilde{\mathbf{x}}_t^\top} \right), & \mathbf{Q}_t &= \begin{bmatrix} (1 - \beta_t) & \beta_t \\ \beta_t & (1 - \beta_t) \end{bmatrix}, \quad \beta_t \in [0, 1], \end{aligned}$$

### III. Theoretical Analysis

#### Energy-guided Sampling for Cross-problem

$$\nabla_x \log p(x|y) = \nabla_x \log p(x) + \nabla_x \log p(y|x)$$

$$\underbrace{\nabla_{\mathbf{x}_t} \log p_{\theta}(\mathbf{x}_t | y_{G'}^*, G')}_{\text{posterior score}} = \underbrace{\nabla_{\mathbf{x}_t} \log p_{\theta}(\mathbf{x}_t | G')}_{\text{pre-trained prior score}} + \underbrace{\nabla_{\mathbf{x}_t} \log p_t(y_{G'}^* | \mathbf{x}_t, G')}_{\text{energy-guided score}}.$$

### III. Theoretical Analysis

#### Energy-guided Sampling for Cross-problem

$$\nabla_x \log p(x|y) = \nabla_x \log p(x) + \nabla_x \log p(y|x)$$

$$\underbrace{\nabla_{\mathbf{x}_t} \log p_{\theta}(\mathbf{x}_t|y_{G'}^*, G')}_{\text{posterior score}} = \underbrace{\nabla_{\mathbf{x}_t} \log p_{\theta}(\mathbf{x}_t|G')}_{\text{pre-trained prior score}} + \underbrace{\nabla_{\mathbf{x}_t} \log p_t(y_{G'}^*|\mathbf{x}_t, G')}_{\text{energy-guided score}}.$$

$$\nabla_{\mathbf{x}_t} \log p_t(y_{G'}^*|\mathbf{x}_t, G') \propto -\nabla_{\mathbf{x}_t} \mathcal{E}(y_{G'}^*, \mathbf{x}_0(\mathbf{x}_t); G'), \quad \mathcal{E}(\cdot|\mathcal{G}) \doteq |y - f(\cdot; G)|$$

### III. Theoretical Analysis

#### Energy-guided Sampling for Cross-problem

$$\nabla_x \log p(x|y) = \nabla_x \log p(x) + \nabla_x \log p(y|x)$$

$$\underbrace{\nabla_{\mathbf{x}_t} \log p_{\theta}(\mathbf{x}_t|y_{G'}^*, G')}_{\text{posterior score}} = \underbrace{\nabla_{\mathbf{x}_t} \log p_{\theta}(\mathbf{x}_t|G')}_{\text{pre-trained prior score}} + \underbrace{\nabla_{\mathbf{x}_t} \log p_t(y_{G'}^*|\mathbf{x}_t, G')}_{\text{energy-guided score}}.$$

$$\nabla_{\mathbf{x}_t} \log p_t(y_{G'}^*|\mathbf{x}_t, G') \propto -\nabla_{\mathbf{x}_t} \mathcal{E}(y_{G'}^*, \mathbf{x}_0(\mathbf{x}_t); G'), \quad \mathcal{E}(\cdot|\mathcal{G}) \doteq |y - f(\cdot; G)|$$

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t, y_{G'}^*, G') \propto p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t, G') p(y_{G'}^*|\mathbf{x}_t, G'),$$

$$p(y_{G'}^*|\mathbf{x}_t, G') = \exp \left( -\frac{1}{\tau} \nabla_{\mathbf{x}_t} f(\tilde{\mathbf{x}}_0(\mathbf{x}_t); G') \right).$$

Reference:

Diffusion Models Beat GANs on Image

<https://arxiv.org/pdf/2105.05233#page=25&zoom=100,144,96>



## IV. Proposed Approach

---

**Algorithm 1:** Energy-guided Diffusion Sampling for Cross-problem Transfer

---

**Input:**

- 1:  $p_\theta$ : Pre-trained diffusion model
- 2:  $G'$ : Target problem instance
- 3:  $T$ : Number of diffusion steps
- 4:  $\tau$ : Energy guidance temperature
- 5:  $K$ : Number of re-inference iterations

**Output:** Optimal solution  $\mathbf{t}_K$  for problem instance  $G'$

- 6: Initialize  $\mathbf{x}_T$  with random binary values;
  - 7: **for**  $k = 1$  **to**  $K$  **do**
  - 8:   Initialize  $\mathbf{x}_T$  with previous best solution  $\mathbf{t}_{k-1}$
  - 9:   **for**  $t = T$  **to** 1 **do**
  - 10:     Compute  $p_\theta(\mathbf{x}_t|G')$  from pre-trained model;
  - 11:     Compute energy gradient:  $\nabla_{\mathbf{x}_t} f(\tilde{\mathbf{x}}_0(\mathbf{x}_t); G')$ ;
  - 12:     Compute  $p_t(y_{G'}^*|\mathbf{x}_t, G') \propto \exp(-\nabla_{\mathbf{x}_t} f/\tau)$ ;
  - 13:     Compute guided posterior:  $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, y_{G'}^*, G') \propto p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, G')p_t(y_{G'}^*|\mathbf{x}_t, G')$ ;
  - 14:     Update next state with Bernoulli Sampling:  
       $\mathbf{x}_{t-1} \sim \text{Cat}(\mathbf{x}_{t-1}; p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, y_{G'}^*, G'))$ ;
  - 15:   **end for**
  - 16:   Decode  $\mathbf{x}_0$  to the best solution  $\mathbf{t}_k$ ;
  - 17: **end for**
  - 18: **return**  $\mathbf{t}_K$
-

## V. Numerical Results

	Method	PCTSP-20		PCTSP-50		PCTSP-100		Avg Gap ↓	Training-based	Training Time ↓
		Gap ↓	Time ↓	Gap ↓	Time ↓	Gap ↓	Time ↓			
OR	Gurobi	<b>0.00%</b>	3.10s	—	—	—	—	—	—	—
	OR-Tools	2.13%	12.31s	4.85%	2.02m	10.33%	5.84m	5.77%	—	—
	ILS (C++)	1.07%	2.13s	<b>0.00%</b>	18.30s	<b>0.00%</b>	56.11s	0.36%	—	—
	ILS (Python 10x)*	63.23%	3.05s	148.05%	4.70s	209.78%	5.27s	140.35%	—	—
Auto-reg*	AM (Greedy)	2.88%	0.02s	17.95%	0.06s	29.24%	0.14s	16.69%	✓	3.5 days
	AM (Sampling)	2.54%	2.43s	14.58%	7.08s	22.20%	15.13s	13.11%	✓	3.5 days
	MDAM (Greedy)	11.76%	41.10s	24.73%	1.31m	30.07%	1.96m	22.19%	✓	4.3 days
	MDAM (Beam Search)	5.88%	2.70m	18.81%	4.77m	26.09%	6.97m	16.93%	✓	4.3 days
	AM (ASP)	12.05%	0.03s	10.34%	0.08s	<b>11.56%</b>	0.18s	11.32%	✓	1.1 days
	AM-FT (Sampling)	<b>1.02%</b>	2.51s	14.11%	8.02s	25.19%	18.21s	13.44%	✓	4.9 days
Diff.	DIFUSCO (TSP)	19.21%	1.04s	18.61%	1.69s	43.42%	2.34s	27.08%	✗	1.5 days
	<b>DIF-guide (Ours)</b>	4.83%	4.58s	<b>7.58%</b>	6.37s	18.1%	10.79s	<b>10.17%</b>	✗	<b>0 days</b>

Table 1: Comprehensive evaluation of cross-scale generalization capabilities across different solver categories on PCTSP instances. Comparison includes exact solvers (Gurobi), OR-based heuristics (OR-Tools, ILS), autoregressive models (AM, MDAM, AM-ASP, AM-FT), and diffusion-based approaches (DIFUSCO, DIF-guide). Performance metrics include optimality gap, inference time, and training time. The best results marked with \* are reported from (Wang et al. 2024). The proposed DIF-guide achieves competitive performance while requiring no training.

## V. Numerical Results

Method	PCTSP-20			PCTSP-50			PCTSP-100		
	Cost	Gap	Time	Cost	Gap	Time	Cost	Gap	Time
DIFUSCO	3.78	19.21%	<b>1.04s</b>	5.20	15.97%	<b>1.35s</b>	8.14	35.61%	<b>2.01s</b>
<b>DIF-Guide (Ours)</b>	<b>3.32</b>	<b>4.83%</b>	5.58s	<b>4.69</b>	<b>3.51%</b>	7.51s	<b>6.67</b>	<b>13.32%</b>	12.20s

Table 2: Zero-shot cross-problem transfer performance comparison between baseline DIFUSCO and our DIF-Guide approach on PCTSP instances. Results show solution cost, optimality gap, and computation time across three problem scales (20, 50, and 100 nodes). The energy-guided sampling uses 50 iterations with greedy decoding strategy.

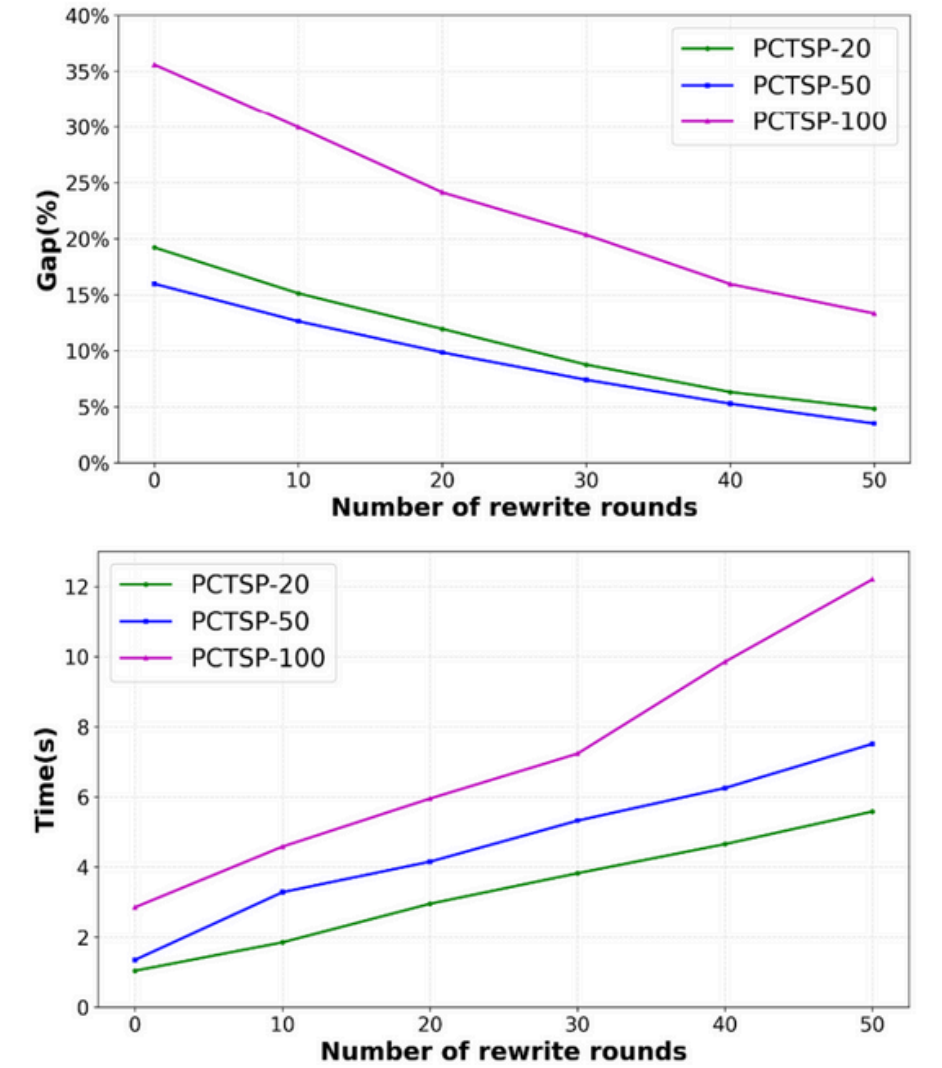


Figure 2: Trade-off between performance and inference time of energy-guided sampling with respect to rewriting rounds.